


# Computer Architecture

## 12. Main Memory

Jianhua Li

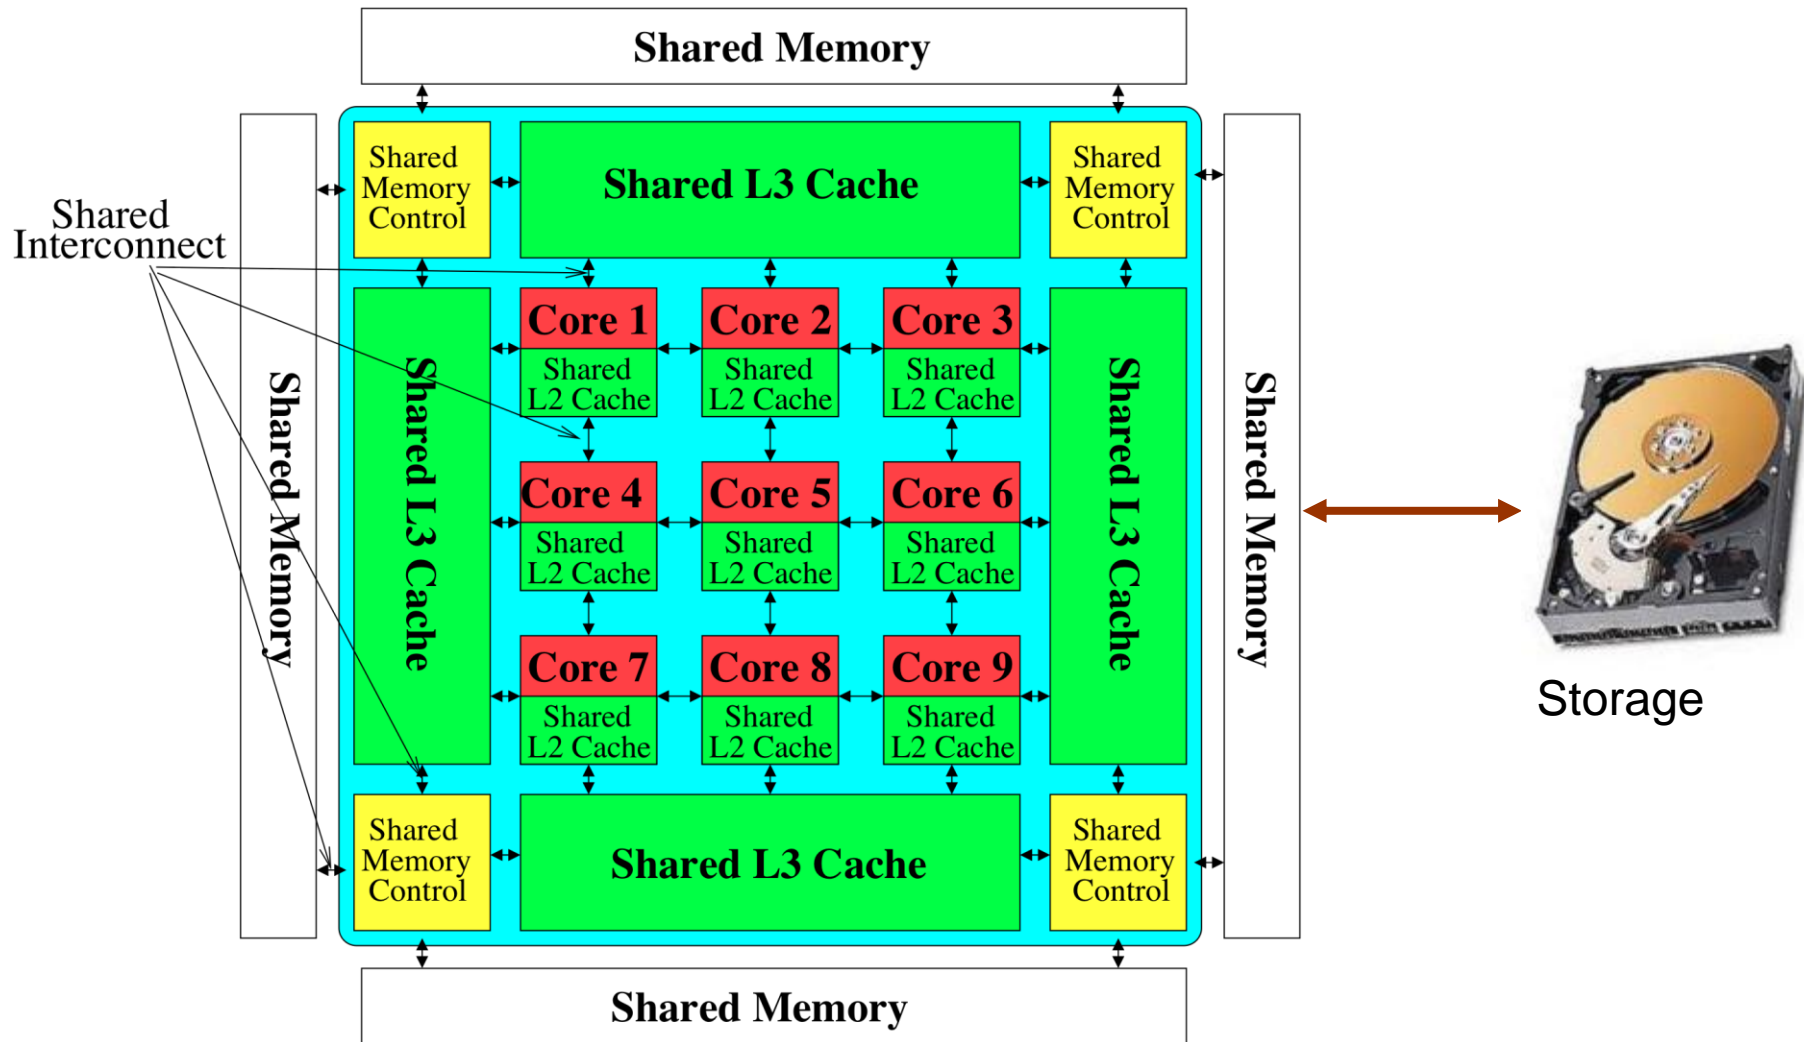
College of Computer and Information  
Hefei University of Technology

# The Main Memory System



- **Main memory is a critical component of all computing systems:** server, mobile, embedded, desktop, sensor
- **Main memory system must scale** (in *size, technology, efficiency, cost, and management algorithms*) to match the growing demands of bandwidths

# Memory System: A Shared Resource View



# State of the Main Memory System


- Recent technology, architecture, and application trends
  - lead to new requirements
  - exacerbate old requirements
- DRAM and memory controllers, as we know them today, are (will be) unlikely to satisfy all requirements
- Some emerging non-volatile memory technologies (e.g., PCM) enable new opportunities: memory + storage merging
- We need to rethink/reinvent the main memory system
  - to fix DRAM issues and enable emerging technologies
  - to satisfy all requirements

# Major Trends Affecting Main Memory


- Need for main memory capacity, bandwidth, QoS increasing
- Main memory energy/power is a key system design concern
- DRAM technology scaling is ending

# Demand for Memory Capacity


- More cores → More concurrency → Larger working set



AMD Barcelona: 4 cores



IBM Power7: 8 cores



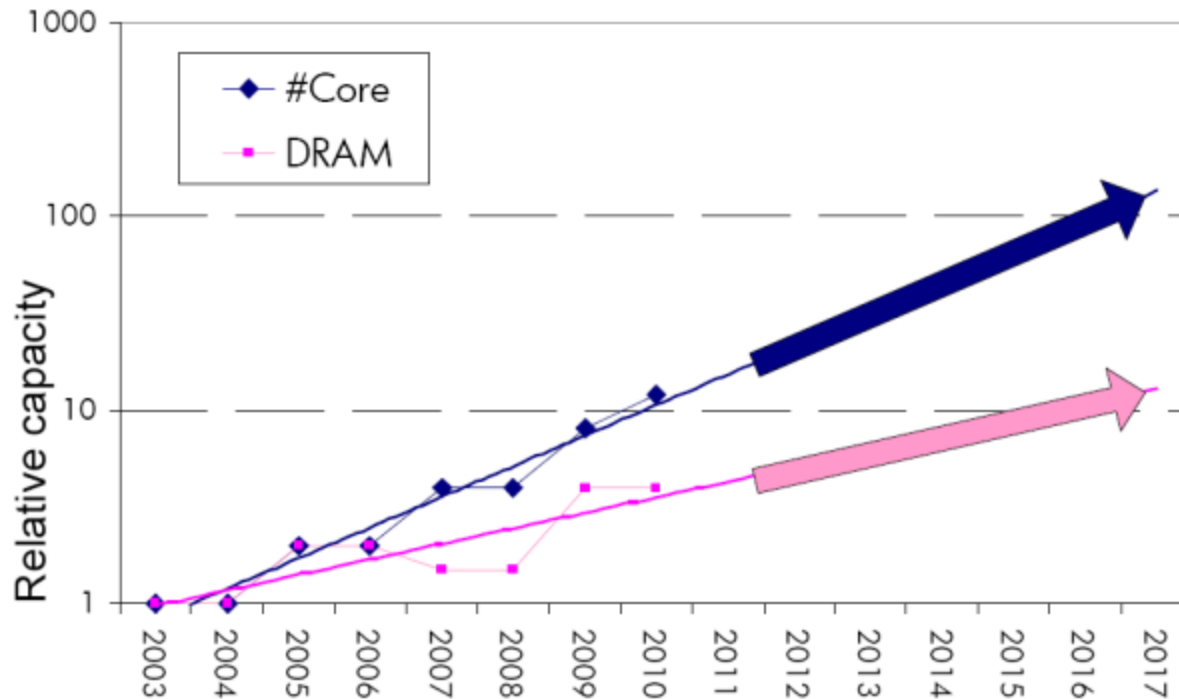
Intel SCC: 48 cores

- Modern applications are (increasingly) data-intensive
- Many applications/virtual machines (will) share main memory
  - Cloud computing/servers: Consolidation to improve efficiency
  - GP-GPUs: Many threads from multiple parallel applications
  - Mobile: Interactive + non-interactive consolidation

# Example: The Memory Capacity Gap

Core count doubling ~ every 2 years

DRAM DIMM capacity doubling ~ every 3 years



- *Memory capacity per core* expected to drop by 30% every two years
- Trends worse for *memory bandwidth per core* !

# Major Trends Affecting Main Memory

- Need for main memory capacity, bandwidth, QoS increasing
  - Multi-core: increasing number of cores
  - Data-intensive applications: increasing demand for data
  - Consolidation: Cloud computing, GPUs, mobile, heterogeneity
- Main memory energy/power is a key system design concern
- DRAM technology scaling is ending



# Major Trends Affecting Main Memory

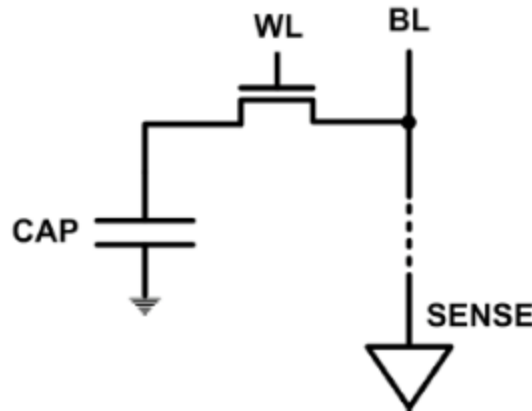
- Need for main memory capacity, bandwidth, QoS increasing
  - **Multi-core**: increasing number of cores
  - **Data-intensive applications**: increasing demand for data
  - **Consolidation**: Cloud computing, GPUs, mobile, heterogeneity
- Main memory energy/power is a key system design concern
  - **IBM servers: ~50% energy spent in off-chip memory hierarchy** [Lefurgy, IEEE Computer 2003]
  - **DRAM consumes power when idle and needs periodic refresh**
- DRAM technology scaling is ending

# Major Trends Affecting Main Memory

- Need for main memory capacity, bandwidth, QoS increasing
  - **Multi-core**: increasing number of cores
  - **Data-intensive applications**: increasing demand for data
  - **Consolidation**: Cloud computing, GPUs, mobile, heterogeneity
- Main memory energy/power is a key system design concern
  - **IBM servers**: ~50% energy spent in off-chip memory hierarchy [Lefurgy, IEEE Computer 2003]
  - **DRAM consumes power when idle and needs periodic refresh**
- DRAM technology scaling is ending
  - ITRS projects **DRAM will not scale easily below X nm**
  - Scaling has provided many benefits:
    - **higher capacity, higher density, lower cost, lower energy**

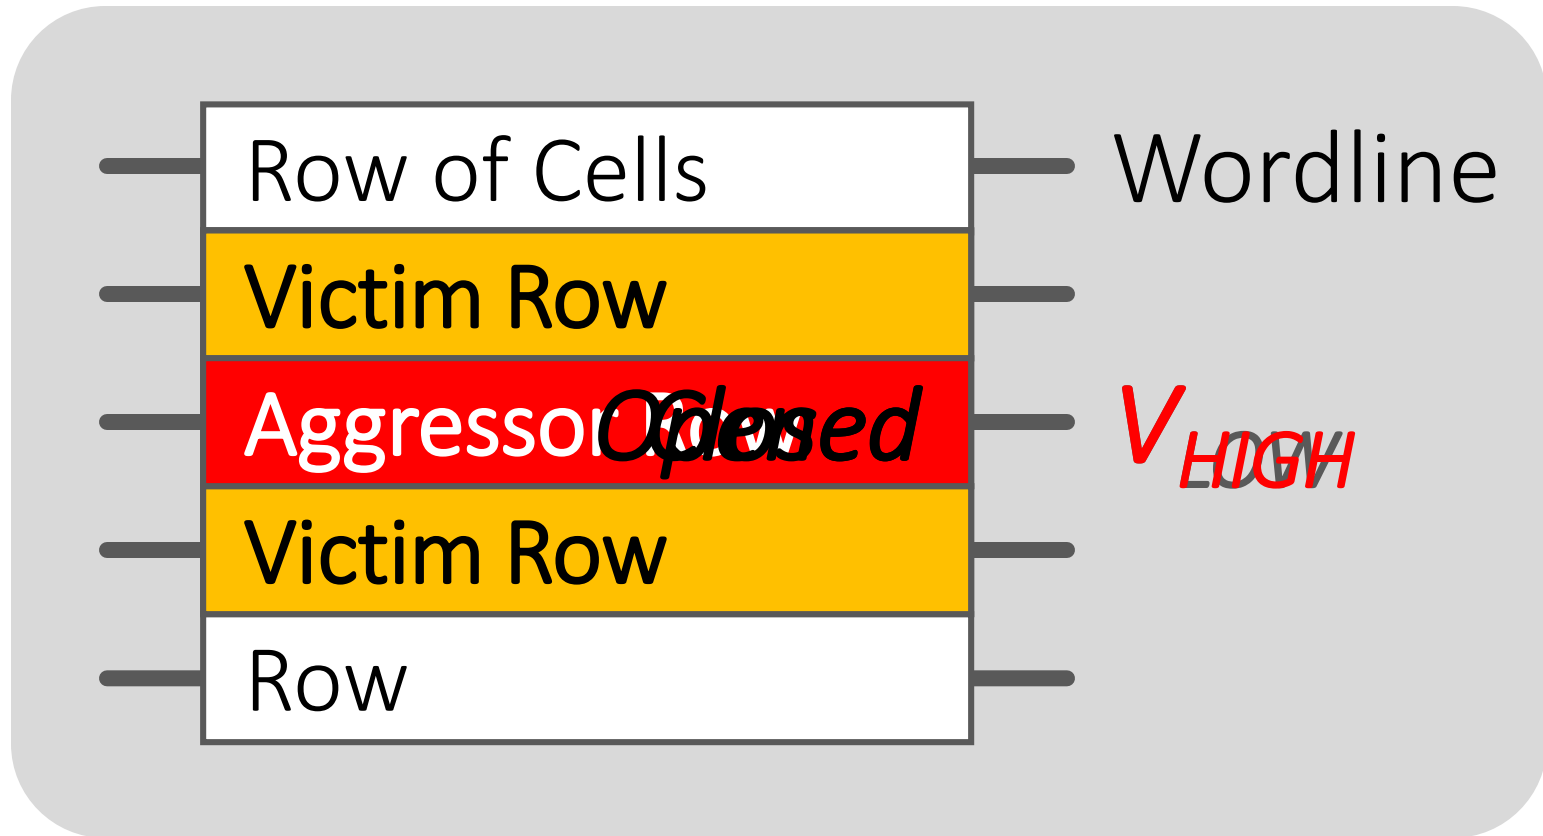
# The DRAM Scaling Problem

- DRAM stores charge in a capacitor (charge-based memory)
  - Capacitor must be large enough for reliable sensing
  - Access transistor should be large enough for low leakage and high retention time
  - Scaling beyond 40-35nm (2013) is challenging [ITRS, 2009]



- DRAM capacity, cost, and energy/power hard to scale


# Evidence of the DRAM Scaling Problem



Repeatedly opening and closing a row enough times within a refresh interval induces **disturbance errors** in adjacent rows in **most real DRAM chips you can buy today**


# Most DRAM Modules Are At Risk

A company




Up to  
 $1.0 \times 10^7$   
errors

B company




Up to  
 $2.7 \times 10^6$   
errors

C company



Up to  
 $3.3 \times 10^5$   
errors


# x86 CPU




# DRAM Module



```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```




# x86 CPU




# DRAM Module



```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```




# x86 CPU



# DRAM Module




```
loop:  
  mov  (X), %eax  
  mov  (Y), %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```






# x86 CPU



# DRAM Module



```
loop:  
  mov  (X),  %eax  
  mov  (Y),  %ebx  
  clflush (X)  
  clflush (Y)  
  mfence  
  jmp  loop
```



# Observed Errors in Real Systems

CPU Architecture	Errors	Access-Rate
Intel Haswell (2013)	22.9K	12.3M/sec
Intel Ivy Bridge (2012)	20.7K	11.7M/sec
Intel Sandy Bridge (2011)	16.1K	11.6M/sec
AMD Piledriver (2012)	59	6.1M/sec

- *A real reliability & security issue*
- *In a more controlled environment, we can induce as many as **ten million** disturbance errors*

# Security Implications

## Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors

*Abstract. Memory isolation is a key property of a reliable and secure computing system — an access to one memory address should not have unintended side effects on data stored in other addresses. However, as DRAM process technology*

## Project Zero

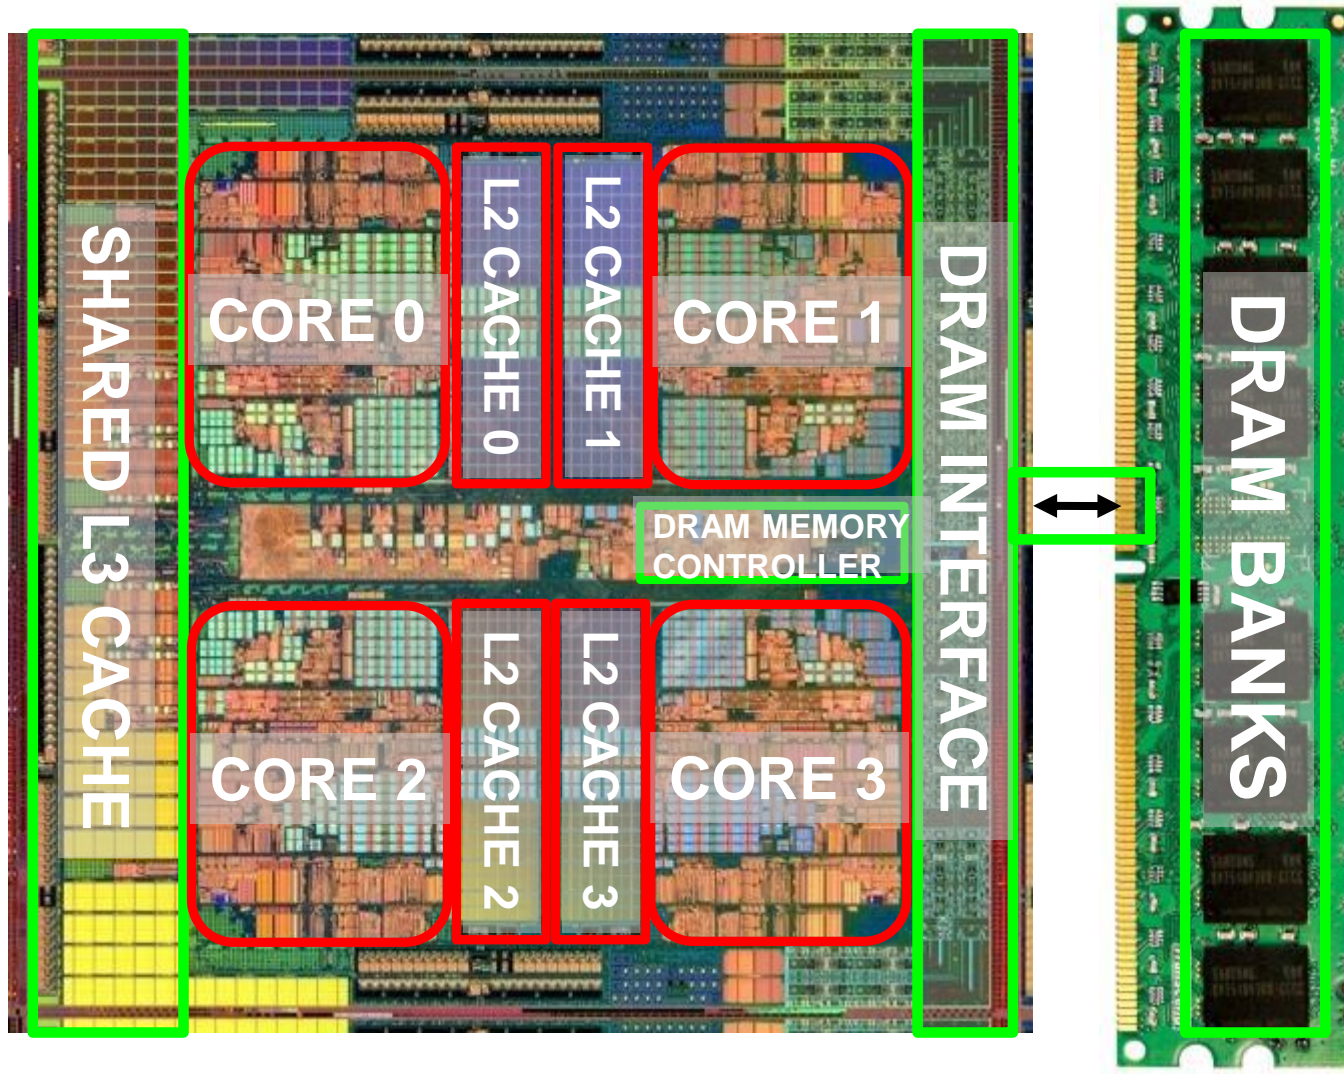
News and updates from the Project Zero team at Google

Monday, March 9, 2015


<http://googleprojectzero.blogspot.com/2015/03/exploiting-dram-rowhammer-bug-to-gain.html>

Exploiting the DRAM rowhammer bug to gain kernel privileges

# Main Memory in the System



# Memory Bank Organization




- Read access sequence:
  1. Decode row address & drive word-lines
  2. Selected bits drive bit-lines
    - Entire row read
  3. Amplify row data
  4. Decode column address & select subset of row
    - Send to output
  5. Precharge bit-lines
    - For next access

# DRAM vs. SRAM

- DRAM
  - Slower access (capacitor)
  - Higher density (1T 1C cell)
  - Lower cost
  - Requires refresh (power, performance, circuitry)
  - Manufacturing requires putting capacitor and logic together
- SRAM
  - Faster access (no capacitor)
  - Lower density (6T cell)
  - Higher cost
  - No need for refresh
  - Manufacturing compatible with logic process (no capacitor)

# DRAM Subsystem Organization

- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column
- Cell



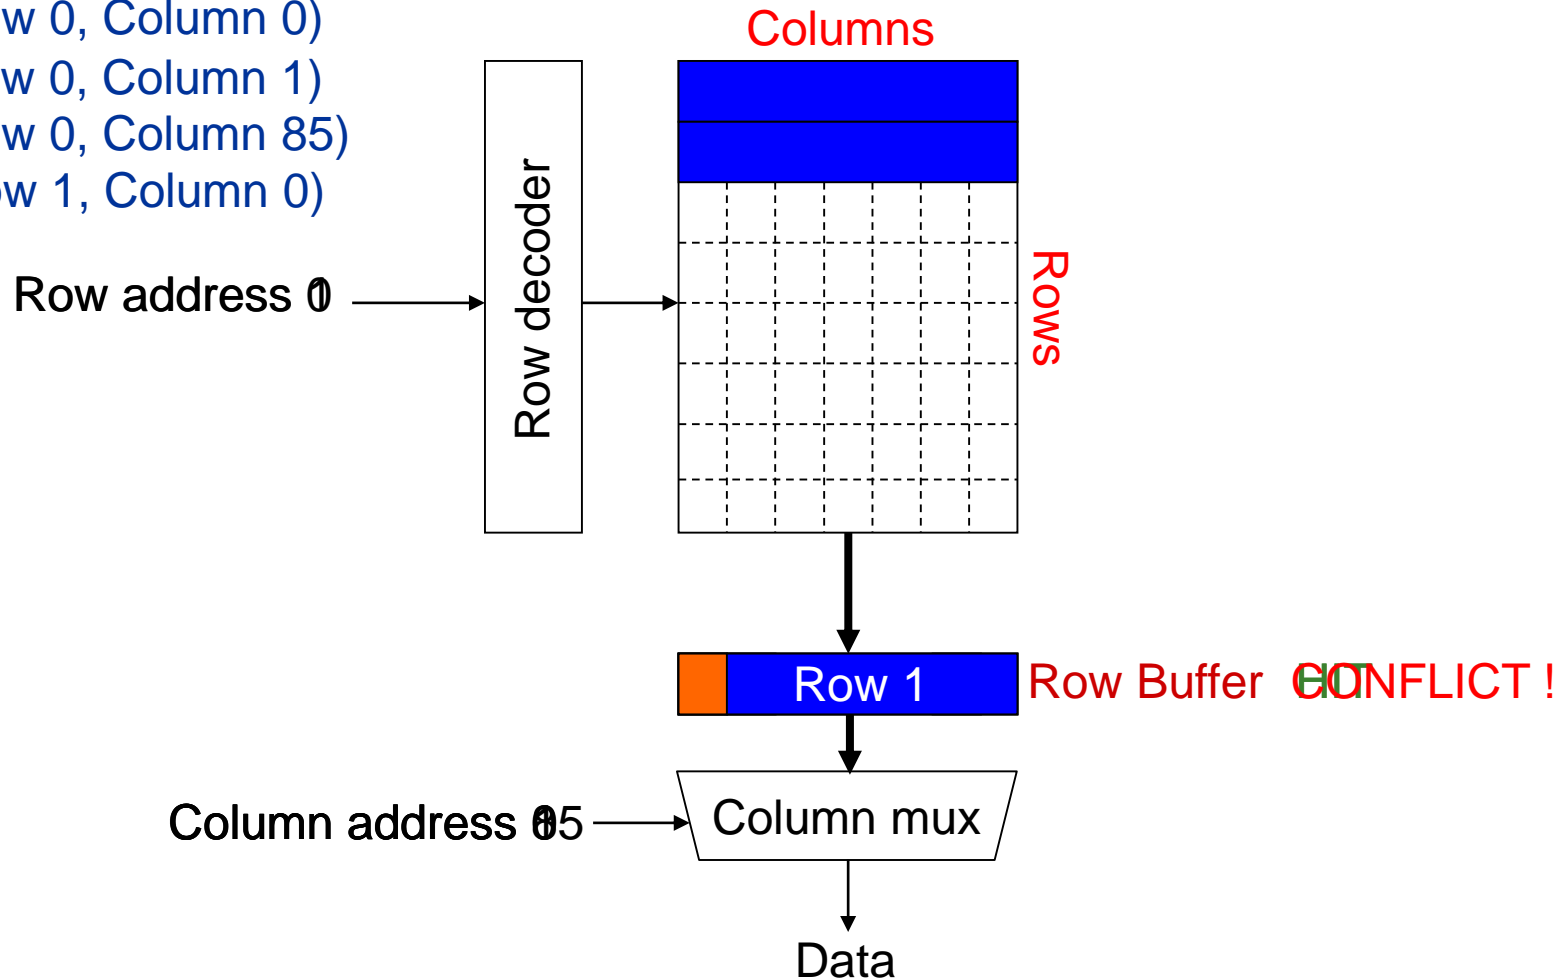
# Page Mode DRAM

- A DRAM bank is a 2D array of cells: rows x columns
- A “DRAM row” is also called a “DRAM page”
- “Sense amplifiers” also called “row buffer”
- Each address is a <row,column> pair
- Access to a “closed row”
  - **Activate** command opens row (placed into row buffer)
  - **Read/write** command reads/writes column in the row buffer
  - **Precharge** command closes the row and prepares the bank for next access
- Access to an “open row”
  - No need for an activate command



# DRAM Bank Operation

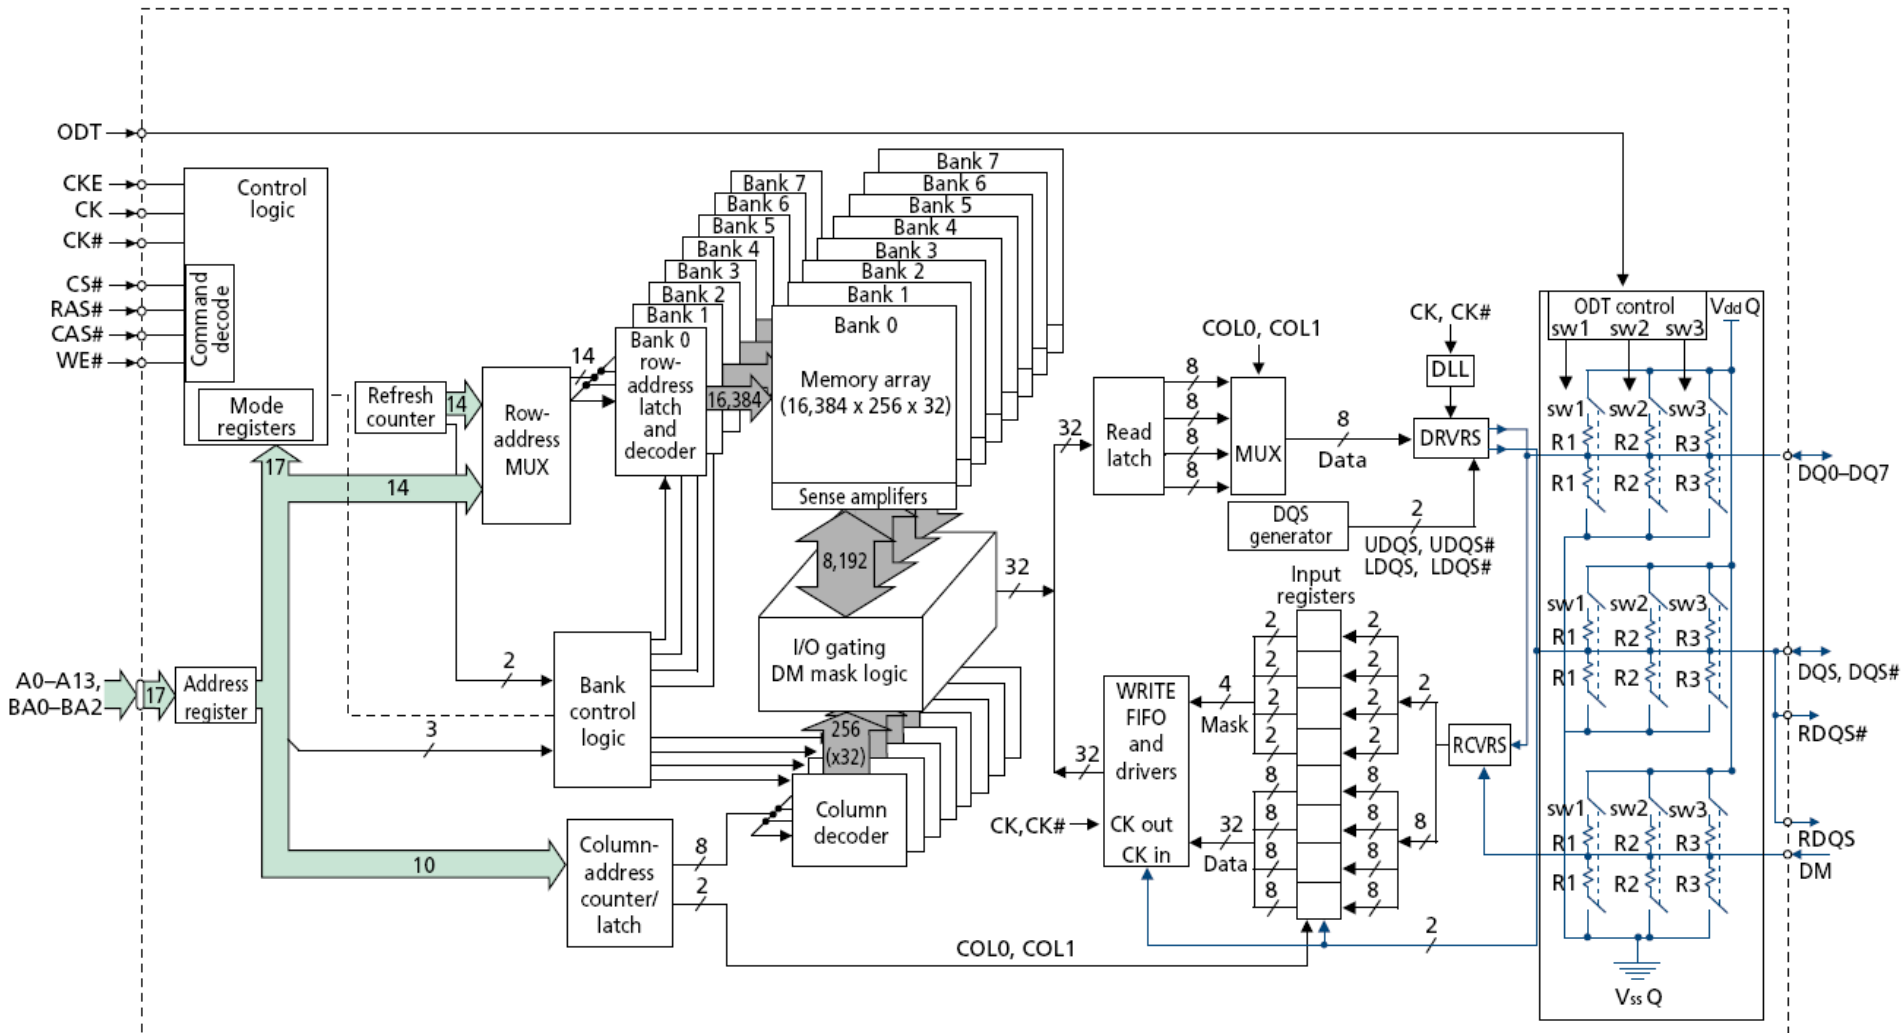
Access Address:  
(Row 0, Column 0)  
(Row 0, Column 1)  
(Row 0, Column 85)  
(Row 1, Column 0)



# The DRAM Chip

- Consists of multiple banks (8 is a common number today)
- Banks share command/address/data buses
- The chip itself has a narrow interface (4-16 bits per read)
- Changing the number of banks, size of the interface (pins), whether or not command/address/data buses are shared has significant impact on DRAM system cost

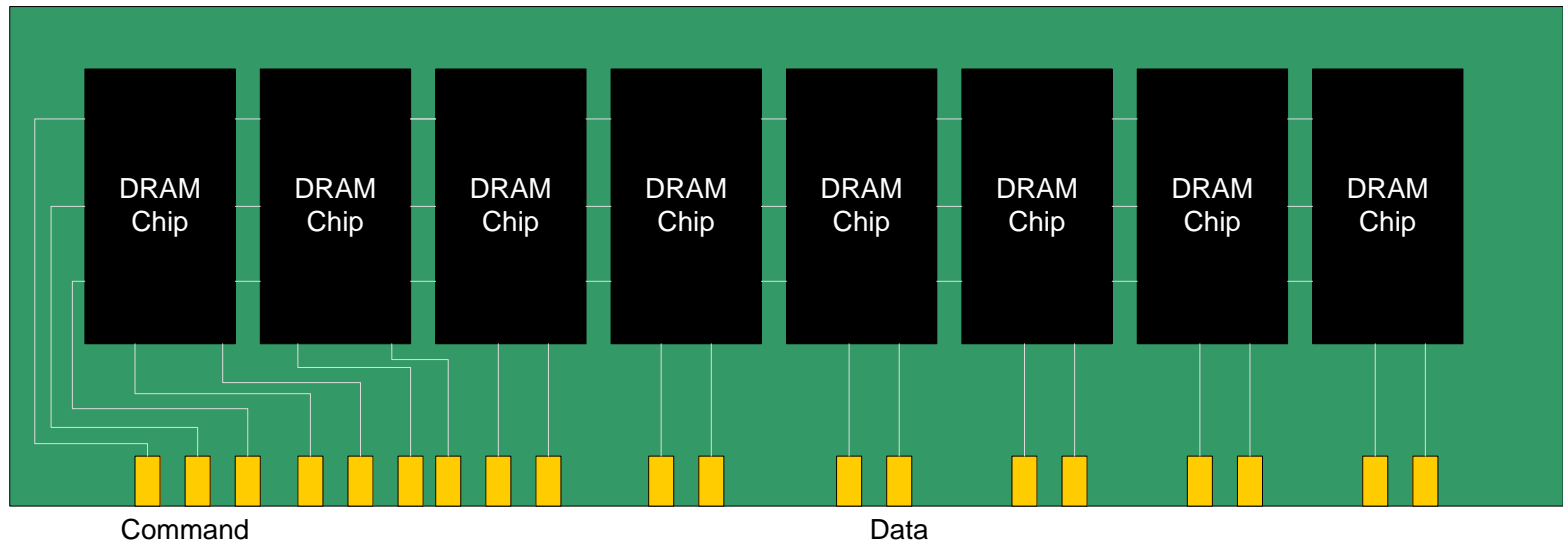
# 128M x 8-bit DRAM Chip



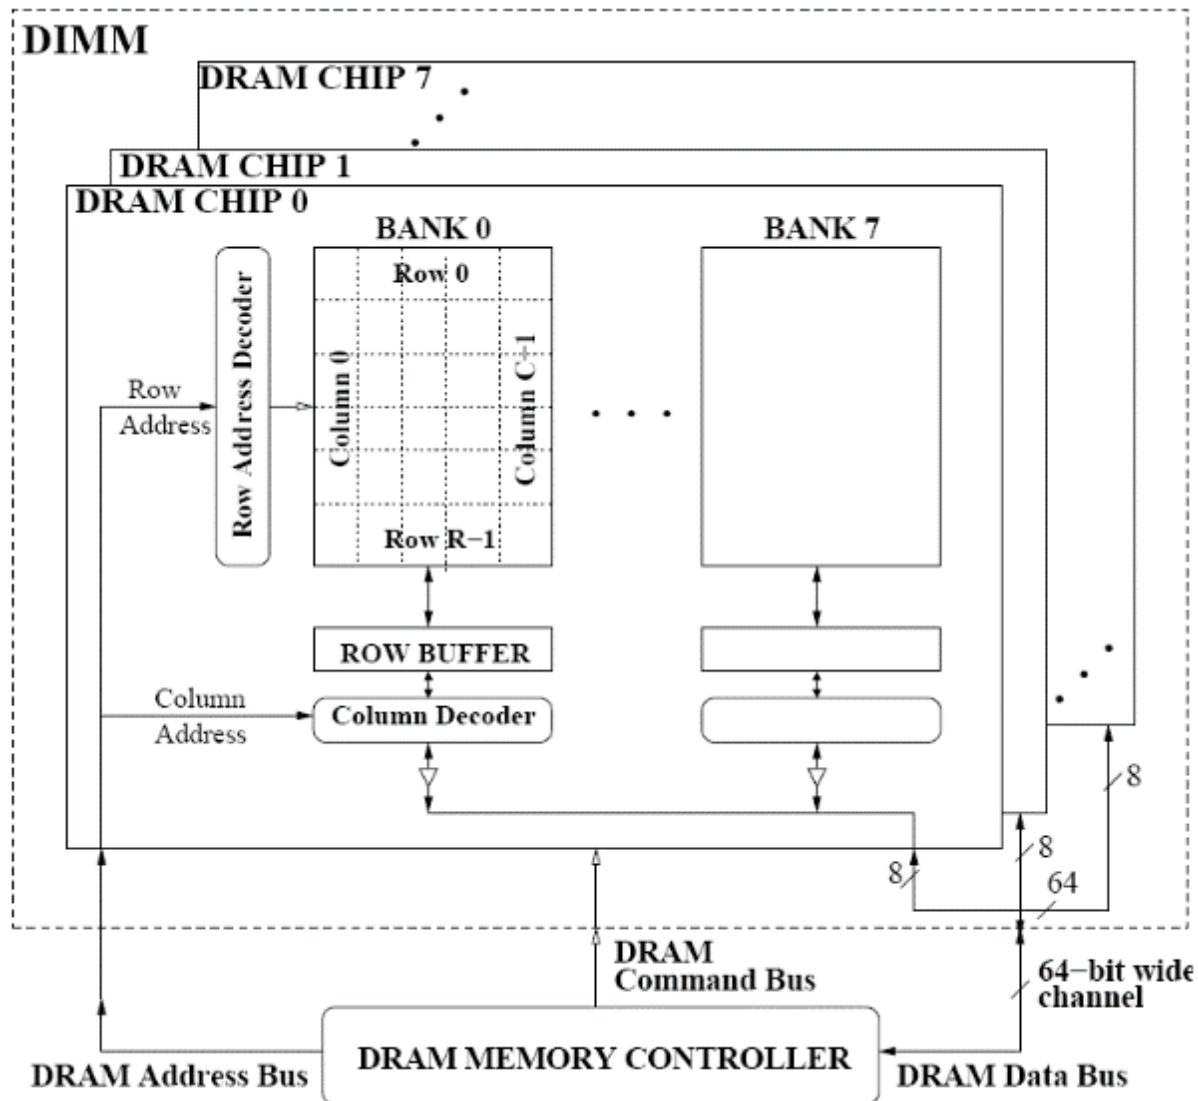
# DRAM Rank and Module

- Rank: Multiple chips operated together to form a wide interface
- All chips comprising a rank are controlled at the same time
  - Respond to a single command
  - **Share address and command buses, but provide different data**
- A DRAM module consists of one or more ranks
  - E.g., DIMM (dual inline memory module)
  - This is what you plug into your motherboard
- If we have chips with 8-bit interface, to read 8 bytes in a single access, use 8 chips in a DIMM

# A 64-bit Wide DIMM (One Rank)



# A 64-bit Wide DIMM (One Rank)



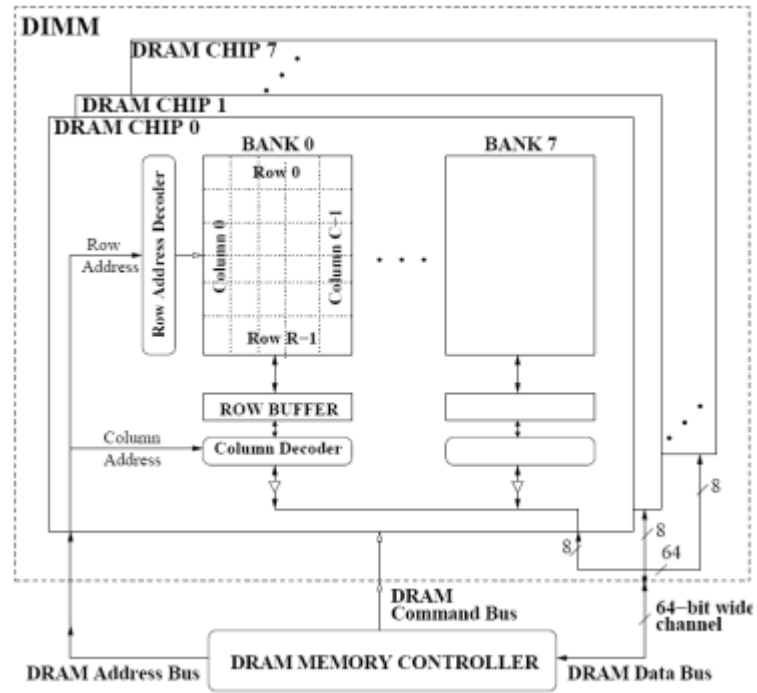
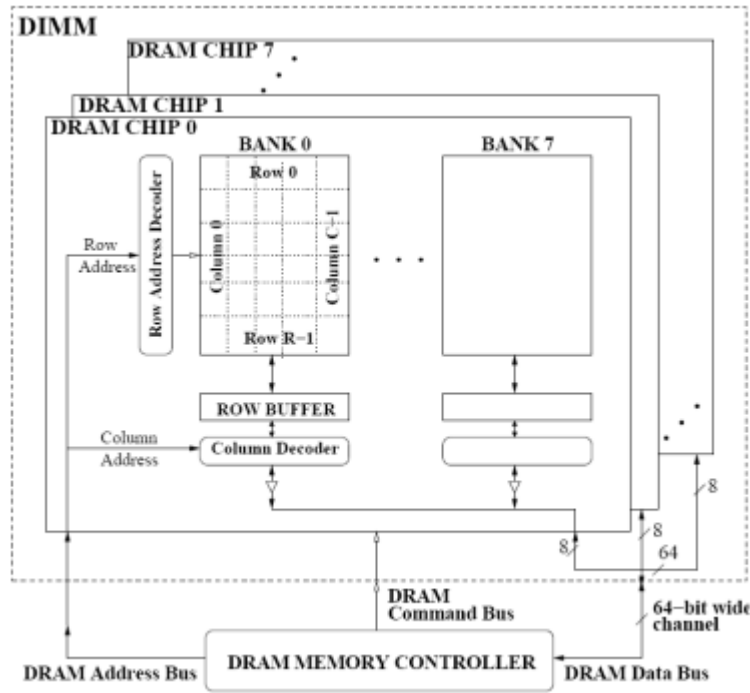
## Advantages:

- Acts like a **high-capacity DRAM chip** with a **wide interface**
- **Flexibility**: memory controller does not need to deal with individual chips

## Disadvantages:


- **Granularity**: Accesses cannot be smaller than the interface width

# DRAM Channels



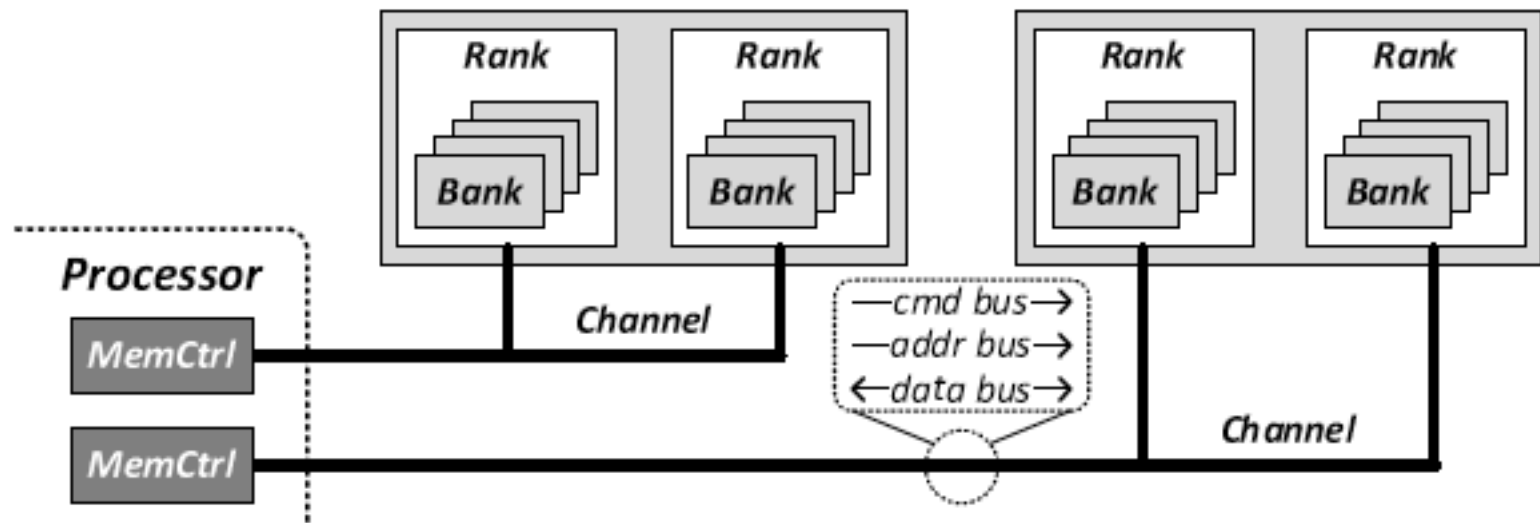
- 2 Independent Channels: 2 Memory Controllers (Above)
- 2 Dependent/Lockstep Channels: 1 Memory Controller with wide interface (Not Shown)

# Generalized Memory Structure






# Generalized Memory Structure



# The DRAM subsystem




# Breaking down a DIMM

**DIMM** (Dual in-line memory module)



Side view



Front of DIMM



Back of DIMM




# Breaking down a DIMM


**DIMM** (Dual in-line memory module)



Side view



Front of DIMM



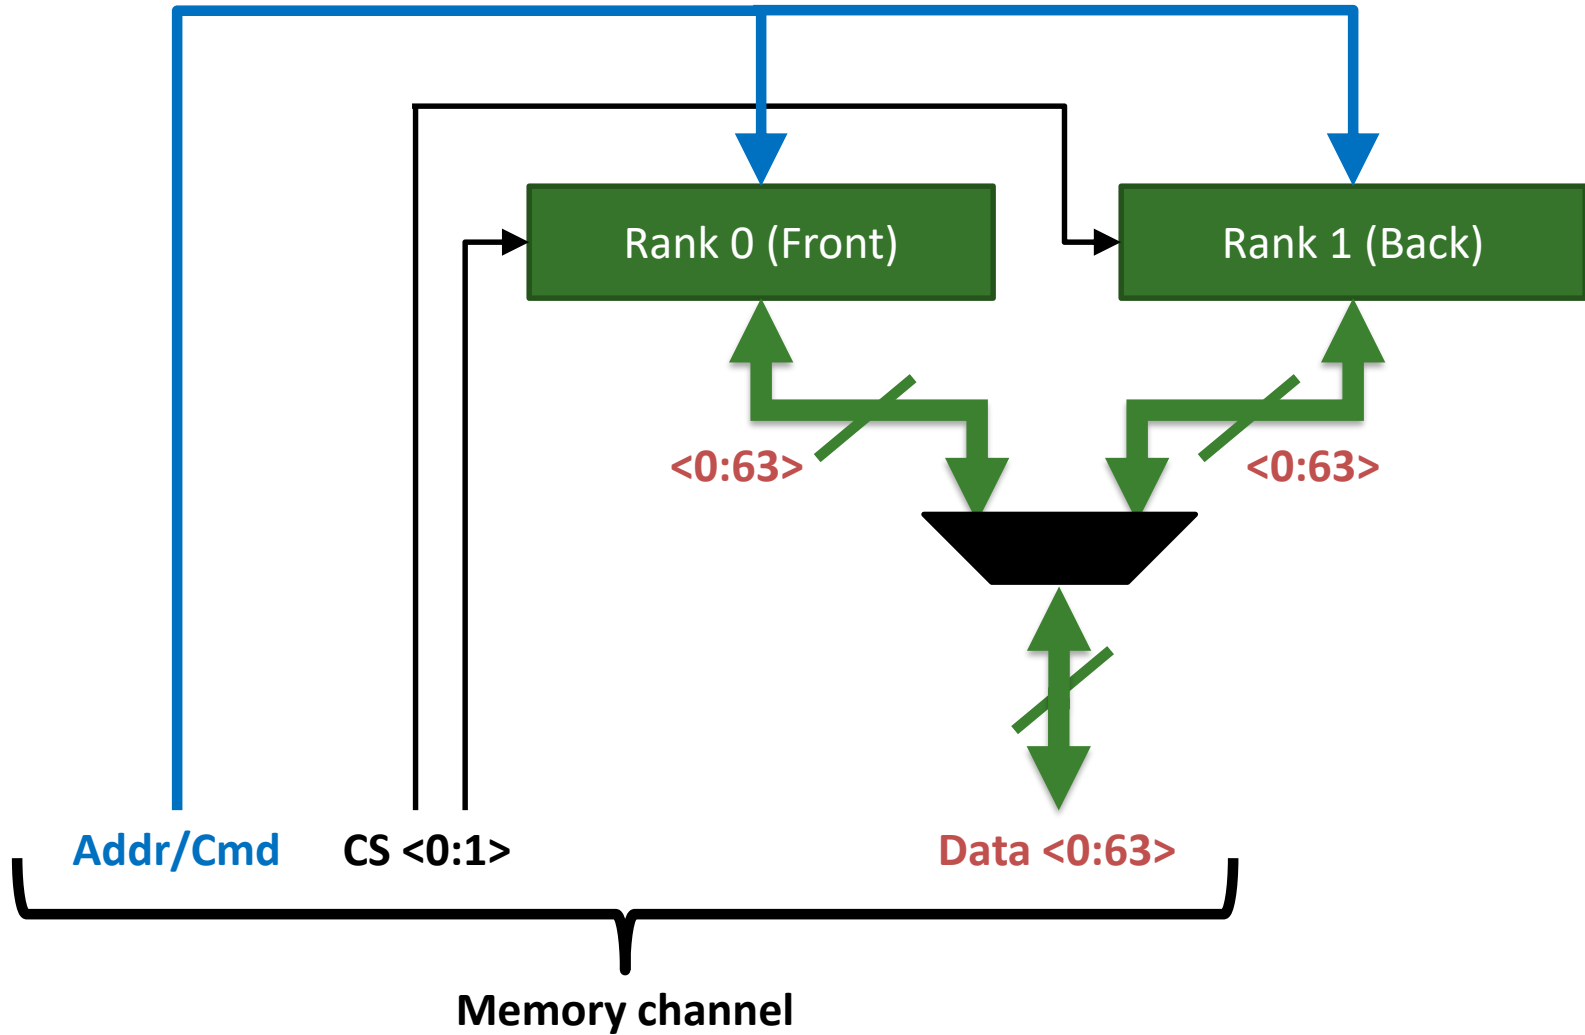
**Rank 0:** collection of 8 chips

Back of DIMM




**Rank 1**


# Rank




# Breaking down a Rank



# Breaking down a Chip




# Breaking down a Bank






# DRAM Subsystem Organization

- Channel
- DIMM
- Rank
- Chip
- Bank
- Row/Column
- Cell




# Example: Transferring a cache block

Physical memory space




# Example: Transferring a cache block

Physical memory space



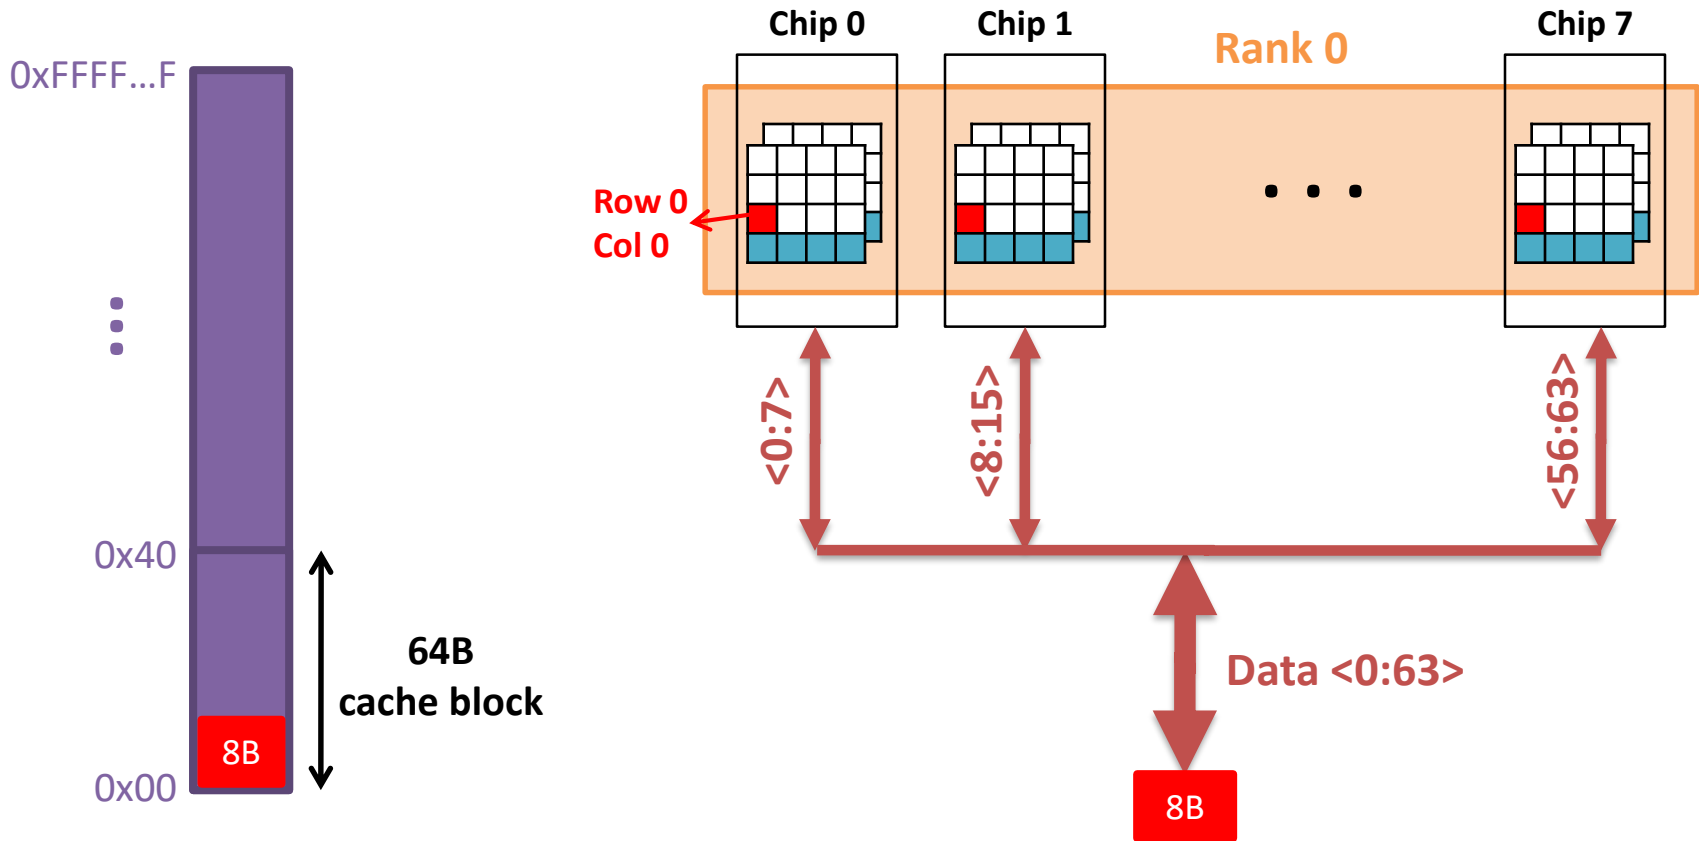
# Example: Transferring a cache block

Physical memory space



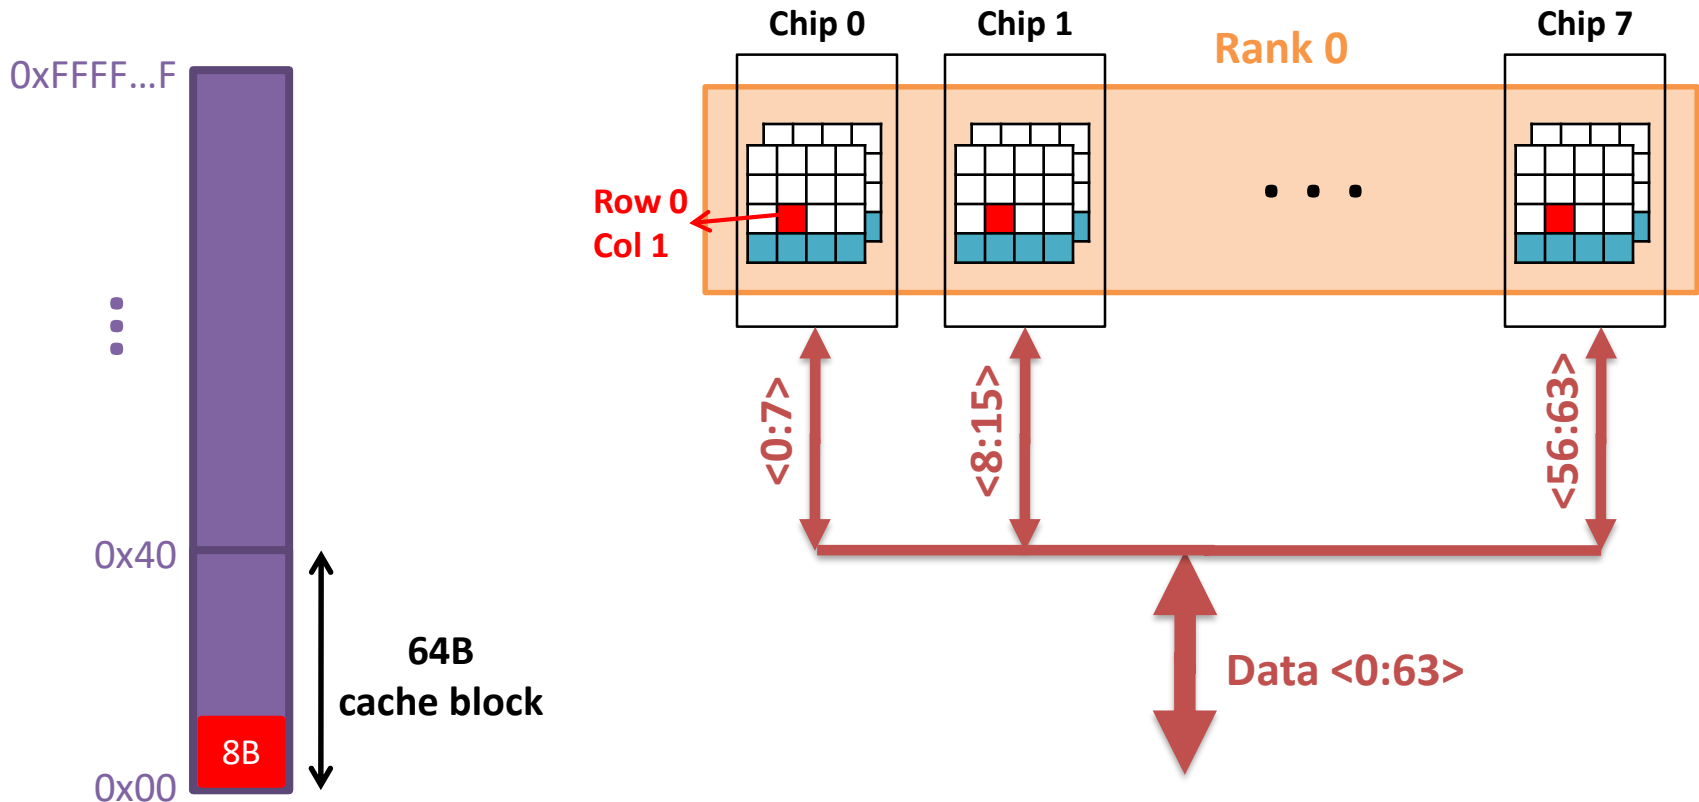
# Example: Transferring a cache block

Physical memory space



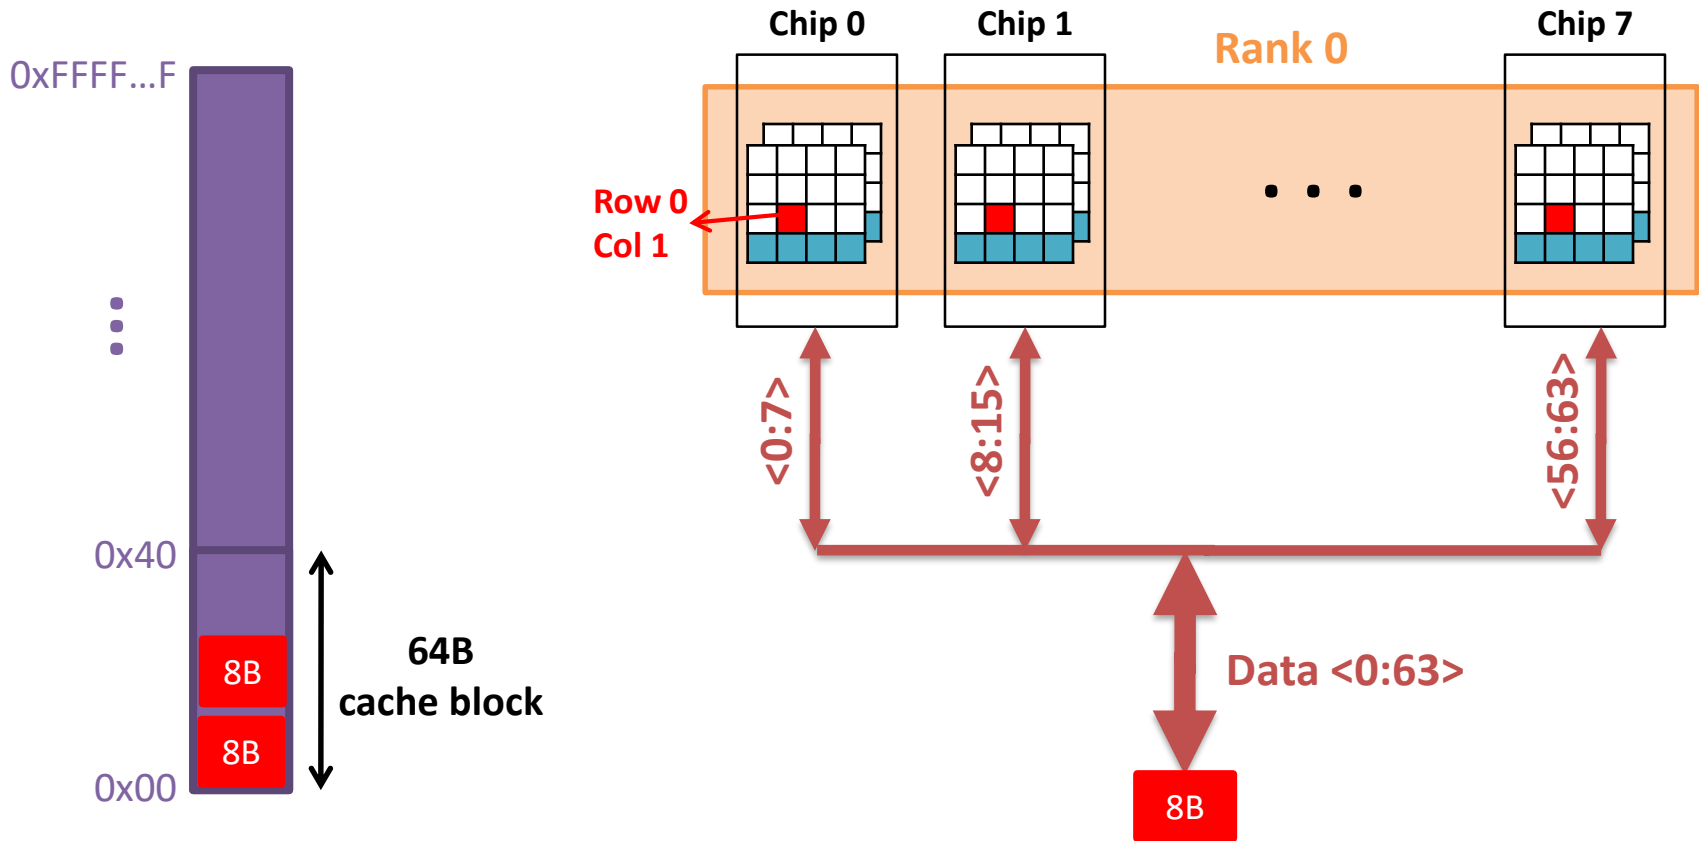
# Example: Transferring a cache block

Physical memory space

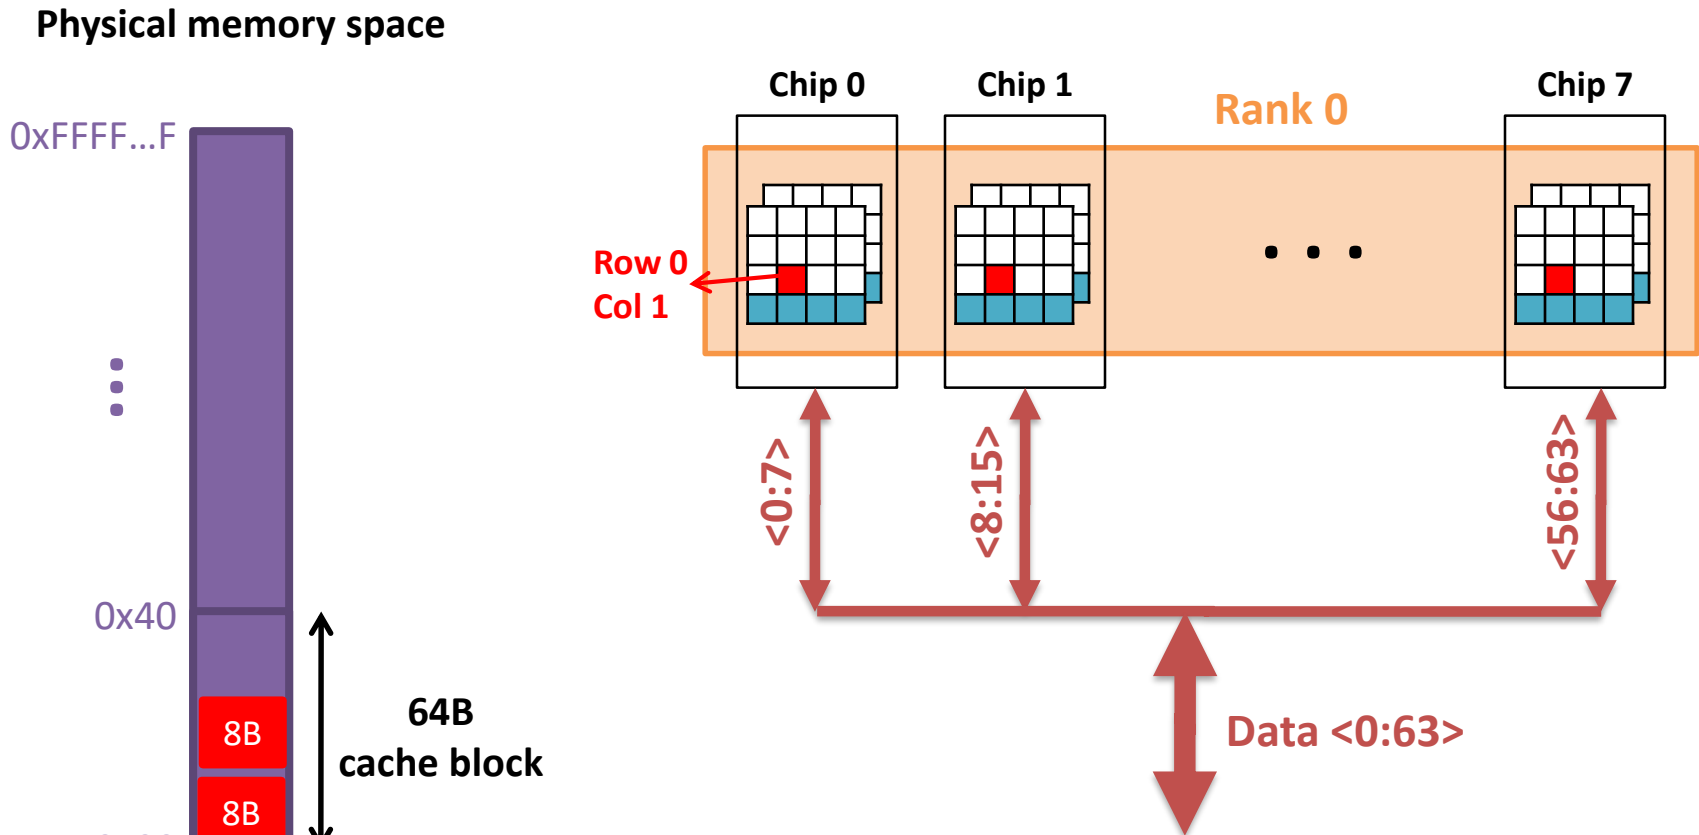


# Example: Transferring a cache block

Physical memory space



# Example: Transferring a cache block



A 64B cache block takes 8 I/O cycles to transfer.

During the process, 8 columns are read sequentially.



# **Next Topic**

## **Virtual Memory**